

Markov chain Monte Carlo and Perfect Simulation

Lecture at Aristotle University of Thessaloniki

Wilfrid S Kendall

Department of Statistics, University of Warwick

15 May 2024



Introduction



Figure 1: Αριστοτέλης 384–322 BCE

Aristotle:

- “Pleasure in the job puts perfection in the work.”
- “The more you know, the more you know you don’t know.”

Handout available on the web: either use the QR-code



or visit <https://wilfridskendall.github.io/talks/Thessaloniki-2024/>.

Sketch of MCMC (I)



Figure 2: Edward Teller (1908-2003)

The original Markov chain Monte Carlo method (**MCMC**) was introduced by Metropolis *et al.* (**1953**). The senior author was Edward Teller (“father of the H-bomb”).

[Fermi once remarked that] Teller was the only monomaniac he knew who had *several* manias (**Brown & May, 2004**).

Sketch of MCMC (II)

- Markov chain basics:

- ▶ Transition probabilities $p(a, b)$ (or transition rates in continuous time: unified view using exponential distribution);
- ▶ Equilibrium probabilities $\pi(a)$, balance, detailed balance $\pi(a)p(a, b) = \pi(b)p(b, a)$, reversibility;
- ▶ Aperiodicity, recurrence (and uniform and geometric recurrence);
- ▶ under detailed balance we can condition by forbidding transitions;

- We can modify any chain, transition probabilities $p(a, b)$, to leave a specified target distribution $\pi(a)$ invariant, by censoring each possible transition $a \rightarrow b$ with probability $\alpha(a, b) \in [0, 1]$ such that

$$\alpha(a, b)\pi(a)p(a, b) = \alpha(b, a)\pi(b)p(b, a);$$

- Common choice: Metropolis-Hastings

$$\alpha(a, b) = \min\{1, (\pi(b)p(b, a))/(\pi(a)p(a, b))\}.$$

- If result still irreducible aperiodic, then $\pi(a)$ is its long-term equilibrium.
- This is MCMC, now of intense interest to statisticians.
- But, physicists always remind us, physicists got there fifty years earlier!

Sketch of MCMC (III)

Given the $\pi(a)$, how to **design** a Markov chain to have this as equilibrium?

- ❶ **Independence sampler**: draw from a fixed probability distribution, apply Metropolis-Hastings;
- ❷ **Random walk Metropolis** or **RWM**: propose move using a random walk, apply Metropolis-Hastings;
- ❸ **Metropolis-adjusted Langevin** or **MALA**: make a Gaussian jump shifted using gradient of $\log \pi$, apply Metropolis-Hastings.

Can mix-and-match! RWM is often favourite: flexible, not too complicated.

Issues:

- Ⓐ **Burn-in**: *How long* till approximate equilibrium?
- Ⓑ **Scaling**: *How big* should be the RWM jump?

Question (B) is about how to get fast mixing. There is a beautiful and useful theory, but that is for another day.

Question (A) is what this lecture is all about.

Sketch of MCMC (IV)

- MCMC practicalities: Burn-in: what to do about it?
 - ▶ Theory tends to be much too pessimistic. Example: Zanella (2015a, 2015b) developed statistical methods for Anglo-Saxon history: a *simplified* model appeared to converge approximately in 10^5 steps (about 1 week on compute cluster), *versus* 10^9 steps in theory (around 2 centuries);
 - ▶ Is (a) one long run better or (b) many short runs? (Option (b) requires starts of short runs spread “evenly” over the sample space — almost as hard in high dimensions as the original problem!)
 - ▶ Diagnostics? (Meta-theorem: for any diagnostic technique there is a chain for which the technique is deceptive!)
 - ▶ Conclusion: effective MCMC requires very careful thought about appropriate length of run — think deeply about the problem!
- Can there ever be a better way?


Perfect Simulation

- The idea of exact simulation / Coupling from the Past (CFTP) / perfect simulation Propp & Wilson (1996) (Persi Diaconis: “Like seeing the landscape of Mars for the first time”);
- Ideas (of “*classic CFTP*”):
 - ▶ extending simulation *backwards* through time,
 - ▶ exploit monotonicity by coupling maximal and minimal processes,
 - ▶ seek coalescence;
- Details for *random-walk-CFTP*, which can be boosted as above to provide simple image reconstruction of an image using Ising model, Propp & Wilson (1996) show how to vary a clever algorithm to get exact samples for **critical** Ising model (this is what impressed Diaconis);
- “Perfect simulation” (WSK, 1998): because everyone knows it isn’t going to be perfect, whereas people might imagine “exact simulation” would somehow miraculously defeat numerical approximation error :-).

An example and some theory

- An intensely visual example, which helps many people see intuitively what is going on here, is *DeadLeaves-CFTP* (WSK & Thönnies, 1999) (technically, *Occluded CFTP*);
- What about cases where monotonicity fails? or there isn't a sensible “maximal” process? WSK (1998):
 - ▶ cross-couple upper and lower envelope processes,
 - ▶ dominate by amenable “dominating process” (time-reversible, can draw from equilibrium, can couple target processes below dominating process);
- Theoretical limits: *in principle*
 - ▶ *Classical CFTP* equivalent to uniform ergodicity (Foss & Tweedie, 1998).
 - ▶ *Dominated CFTP* is achievable under geometric ergodicity (WSK, 2004).
 - ▶ It is even possible to carry out Dominated CFTP in some **non**-geometrically ergodicity cases [Connor & WSK (2007); *nb* corrigendum];
- We can use *Dominated CFTP* to carry out perfect simulation for stable point processes (WSK & Møller, 2000);
- Detailed expositions are given by WSK (2005), Huber (2015). WSK (2015) shows how to implement CFTP in R.

Applications to Queues and Epidemics



<https://covidposters.github.io/>

Figure 3: An illustration introducing *both* queues *and* epidemics!

Perfect Queues

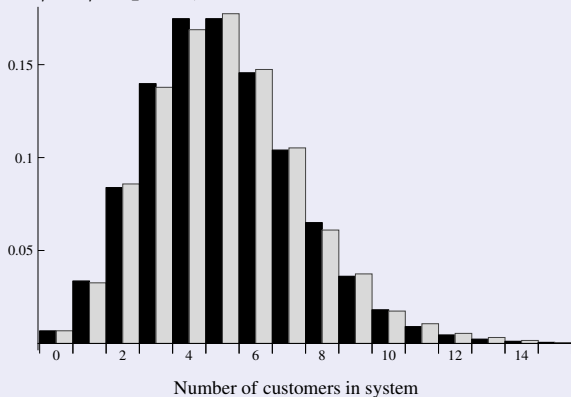
The simplest queuing model ($M/M/1$: Poisson arrivals, exponential service times, single server) can be analyzed very thoroughly indeed! However:

- Poisson arrivals are not unreasonable, but exponential service times are ludicrous. The $M/G/1$ case (general service time *for just one server*) can be analyzed using the “embedded chain” (sample at each departure);
- Multi-server case: computation of *eg* waiting-time distribution is out of reach so use simulation (and insights from [Kiefer & Wolfowitz, 1955](#));
- Sigman ([2011](#)) shows how to do CFTP in the “super-stable” case (traffic so low that it could have been handled by just one server), using Dominated CFTP and comparing to a “Processor-Sharing” discipline.

- Connor & WSK (2015) extend Sigman (2011), showing how to apply Dominated CFTP to simulate (sub-critical!) queues perfectly; now generalized by others to the case of non-Poissonian inter-arrival times. (Technical point: pathwise domination needs service times to be assigned in order of commencement of service!) The idea is
 - ▶ dominate $M/G/c$ FCFS (FCFS means first come first served) by $M/G/c$ RA = $[M/G/1 \text{ RA}]^c$ (RA means assign to individual servers on arrival);
 - ▶ use fact that $M/G/1$ FCFS and $M/G/1$ PS *workloads* agree (PS means Processor Sharing: pool servers to serve everyone simultaneously) and $M/G/1$ PS can be simulated backwards in time;
 - ▶ so $[M/G/1 \text{ PS}]^c$ can be used to provide Dominated CFTP.
- Connor & WSK (2015) describe
 - (a) CFTP coupling when dominating process empties,
 - (b) and a faster CFTP coupling using upper and lower processes starting respectively at dominating process and at empty state.

Results (I)


Histogram of customer numbers for $M/M/c$ queue in equilibrium: arrival rate 10, service rate 2, and 10 servers, comparing theory (available for $M/M/c$ queue) with results of Connor & WSK (2015) algorithm.



Results (II)


Comparison of log-run times for

- Ⓐ CFTP coupling when dominating process empties (solid bars),
- Ⓑ a faster CFTP coupling using upper and lower processes (grey bars).



Results (III)

Marginal distributions of last six coordinates of the equilibrium Kiefer-Wolfowitz workload vector (anticipated work in system) for an $M/G/c = 25$ queue with Uniform(0, 1) service times and arrival rate 25.



Perfect Epidemics

S-I-R deterministic epidemic: susceptibles s , infectives i , removals r :

$$\begin{aligned}s' &= -\alpha i s, \\ i' &= (\alpha s - \beta) i, \\ r' &= \beta i.\end{aligned}$$

(Total population $s + i + r = n$ is constant.)

S-I-R stochastic epidemic: a Markov chain (S, I, R) with transitions

Infection: $S \rightarrow S - 1, \quad I \rightarrow I + 1$ at rate $\alpha I S$,

Removal: $I \rightarrow I - 1, \quad R \rightarrow R + 1$ at rate βI .

Both models assume homogeneous mixing.

The first question asked about a new epidemic

“What is the R-number?”

The R-number is $\alpha n/\beta$: mean number of new infectives produced per infective at *start* of epidemic.

Whittle’s threshold theorem: R-number $\gg 1$ means positive chance of epidemic infecting significant proportion of the population.

Wikipedia: “The British-registered *Diamond Princess* was the first cruise ship to have a major [COVID-19] outbreak on board, with the ship quarantined at Yokohama from 4 February 2020 for about a month. Of 3711 passengers and crew, around 700 people became infected and 9 people died.”

Evidently $\alpha n/\beta \gg 1$ – as was sadly later confirmed, a sorrow for us all.



Inference on the R-number

Important, because the R-number controls severity of epidemic. However:

- 1 It's **tough**. *Either* massive assumptions (homogeneous mixing) or far too many parameters;
- 2 It's **really tough**. It's hard to get good information about infection times;
- 3 It's **especially tough** early on. You most want to know the answer when there is hardly any information available (I devised a simplified exercise for a Warwick second-year statistics module to show how tough this can be);
- 4 Markov chain Monte Carlo (MCMC) can be used (O'Neill & Roberts, 1999) but what about burn-in?
- 5 Can we use perfect simulation?

An easier question


- ① Suppose we know n , α , β , observe removals, but *don't* observe infections which must be inferred.
- ② We need a new chain: namely, the whole S-I-R trajectory evolving in *algorithmic time* using varying pattern of potential infections and removals.
- ③ Visualize n time lines, along which are scattered incidents:
 - ▶ potential removals, activated if time line is infected;
 - ▶ potential infections, activated if time line is infected *and* if designated target time line is lowest uninfected time line.
- ④ Using Poisson point processes of appropriate rates to scatter these incidents, we obtain an S-I-R epidemic.
- ⑤ Now let the point patterns evolve in *algorithmic time*, adding and removing incidents according to spatial immigration-death processes.
- ⑥ Result is a trajectory-valued chain which has unconditioned S-I-R as equilibrium.

Conditioning on observed removals

- The trajectory-valued chain is *reversible*.
- So if we forbid the evolution to get rid of observed removals, and forbid it to introduce new activated removals, then the modified chain has invariant probability measure which conditions on observed pattern of removals. Implications:
 - ▶ conditioned removals can change time line (if still activated) but not time of occurrence;
 - ▶ removals can be introduced only if they don't activate;
 - ▶ sometimes infections cannot be removed (because that would result in losing a conditioned removal).
- More housekeeping details required to define algorithm precisely and make sure reversibility and monotonicity still work.
- Need to ensure irreducibility (or otherwise equilibrium will depend on starting point).
- Does this produce a feasible algorithm?

Example

- Smallpox outbreak in a closed community of 120 individuals in Abakaliki, Nigeria (Bailey, 1975; O'Neill & Roberts, 1999).
- **Assume**
 - ▶ first observed removal is also the first removal: under a plausible improper prior we can then deduce what is the distribution of infectives I_0 at time 0;
 - ▶ all removals are recorded;
 - ▶ there are no further removals after the last observed removal.
- Coding in *julia* (Bezanson *et al.*, 2017), we can derive a perfect simulation drawing from conditional distribution of pattern of infections.



So what?

- You may be wondering, why this emphasis on unobserved infections given fixed α and β , when what we really want is inference on R-number $\alpha n/\beta$ for *unknown* α and β ?
- Good question. But a re-weighting argument allows us to get (unbiased) estimates based on *different* α and β . Essentially the perfect simulation provides an exact simulation-based computation which permits us to integrate out the pattern of unobserved infections.
- This means we can (**work in progress**, Connor & WSK, 2024)
 - ▶ construct a steepest ascent algorithm (in effect, a variant on a Robbins-Monro stochastic optimization algorithm) to find *maximum a posteriori* estimates of α and β ;
 - ▶ or even, with some computational effort, compute the entire posterior joint density for α and β !

Conclusion

- You may not need to put up with burn-in issues when doing MCMC. Consider the possibility of using perfect simulation!
- CFTP works even for significantly complex and relevant models of real-life phenomena;
- *Of course* really detailed models are still going to resist perfect simulation: but it will always be helpful to compare with a simpler model (using fewer parameters!).
- CFTP is clearly an important tool to be considered by the investigator seeking to implement accurate and informative MCMC.

References I

- Bailey, N.T.J. (1975) *The mathematical theory of infectious diseases and its applications*, 2nd Ed. ed. Griffin.
- Bezanson, J., Edelman, A., Karpinski, S., & Shah, V.B. (2017) Julia: A Fresh Approach to Numerical Computing. *SIAM Review*, **59**, 65–98.
- Brown, H. & May, M. (2004) Edward Teller in the Public Arena. *Physics Today*, **57**, 51–53.
- Connor, S.B. & WSK (2007) Perfect simulation for a class of positive recurrent Markov chains. *Annals of Applied Probability*, **17**, 781–808.
- Connor, S.B. & WSK (2015) Perfect simulation of M/G/c queues. *Advances in Applied Probability*, **47**, 1039–1063.
- Connor, S.B. & WSK (2024) Perfect Epidemics.
- Foss, S.G. & Tweedie, R.L. (1998) Perfect simulation and backward coupling. *Stochastic Models*, **14**, 187–203.
- Fraser, C. & Others (2023) OpenABM-Covid19: Agent-based model for modelling the Covid-19 and Contact-Tracing.
- Huber, M.L. (2015) *Perfect Simulation*. Boca Raton: Chapman; Hall/CRC.
- Kiefer, J. & Wolfowitz, J. (1955) On the Theory of Queues With Many Servers. *Transactions of the American Mathematical Society*, **78**, 1.

References II

- Metropolis, N., Rosenbluth, A.W., Rosenbluth, M.N., Teller, A.H., & Teller, E. (1953) Equation of State Calculations by Fast Computing Machines. *The Journal of Chemical Physics*, **21**, 1087.
- Mizumoto, K., Kagaya, K., Zarebski, A., & Chowell, G. (2020) Estimating the asymptomatic proportion of coronavirus disease 2019 (COVID-19) cases on board the Diamond Princess cruise ship, Yokohama, Japan, 2020. *Eurosurveillance*, **25**, 6pp.
- O'Neill, P.D. & Roberts, G.O. (1999) Bayesian Inference for Partially Observed Stochastic Epidemics. *Journal of the Royal Statistical Society Series A: Statistics in Society*, **162**, 121–129.
- Propp, J.G. & Wilson, D.B. (1996) Exact sampling with coupled Markov chains and applications to statistical mechanics. *Random Structures and Algorithms*, **9**, 223–252.
- R Development Core Team (2010) R: A Language and Environment for Statistical Computing.
- Sigman, K. (2011) Exact simulation of the stationary distribution of the FIFO M/G/c queue. *Journal of Applied Probability*, **48**, 209–213.
- WSK (1998) Perfect Simulation for the Area-Interaction Point Process. *Probability towards 2000* (Accardi, L. & Heyde, C.C. eds). Springer-Verlag, pp. 218–234.
- WSK (2004) Geometric ergodicity and perfect simulation. *Electronic Communications in Probability*, **9**, 140–151.
- WSK (2005) Notes on Perfect Simulation. Singapore: World Scientific, pp. 93–146.

References III

- WSK (2015) Introduction to CFTP using R. *Stochastic geometry, spatial statistics and random fields, Lecture notes in mathematics*. Springer, pp. 405–439.
- WSK & Møller, J. (2000) Perfect simulation using dominating processes on ordered spaces, with application to locally stable point processes. *Advances in Applied Probability*, **32**, 844–865.
- WSK & Thönnies, E. (1999) Perfect simulation in stochastic geometry. *Pattern Recognition*, **32**, 1569–1586.
- Zanella, G. (2015b) Bayesian Complementary Clustering, MCMC, and Anglo-Saxon Placenames (PhD Thesis).
- Zanella, G. (2015a) Random partition models and complementary clustering of Anglo-Saxon place-names. *Annals of Applied Statistics*, **9**, 1792–1822.

Technical information

Image	Attribution	
Aristotle	After Lysippos	<i>Public domain</i> via Wikimedia Commons
Edward Teller	Lawrence Livermore National Laboratory restored by w>User:Greg L, Papa Lima Whiskey	CC BY-SA 3.0 via Wikimedia Commons
Perfect Ising	Result of code written by WSK	
Dead leaves	Result of code written by WSK	
Queues	https://covidposters.github.io/	<i>Open source</i>
$M/M/c$ customers	Result of code written by Stephen Connor	
$M/M/c$ runtimes	Result of code written by Stephen Connor	
$M/M/c$ loads	Result of code written by Stephen Connor	
Diamond Princess	Alpsdake	CC BY-SA 4.0
Epidemic	Result of code written by WSK	

These notes were produced from Thessaloniki-2024.qmd:

git repository information.

Version: 0.0.46 (Wed May 8 16:11:07 2024 +0100)

Author: Wilfrid Kendall <W.S.Kendall@warwick.ac.uk>

Date: Tue Wed 8 17:55:14 2024 +0100
